

Computational Photography

Week 7

Instructor: Lou Kratz

Image Warping

All 2D Linear Transformations

- Linear transformations are combinations of ...

- Scale,
 - Rotation,
 - Shear, and
 - Mirror

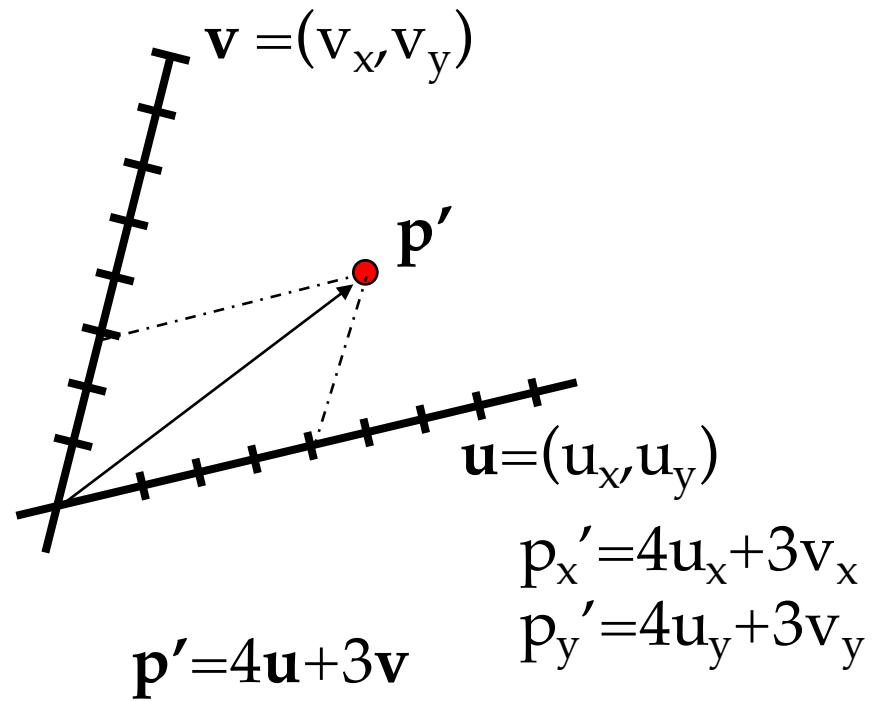
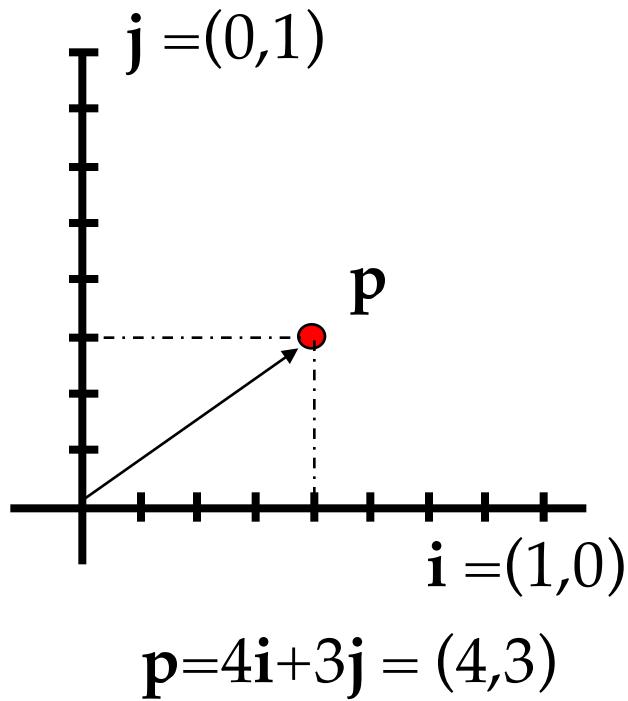
$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

- Properties of linear transformations:

- Origin maps to origin
 - Lines map to lines
 - Parallel lines remain parallel
 - Ratios are preserved
 - Closed under composition

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} e & f \\ g & h \end{bmatrix} \begin{bmatrix} i & j \\ k & l \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

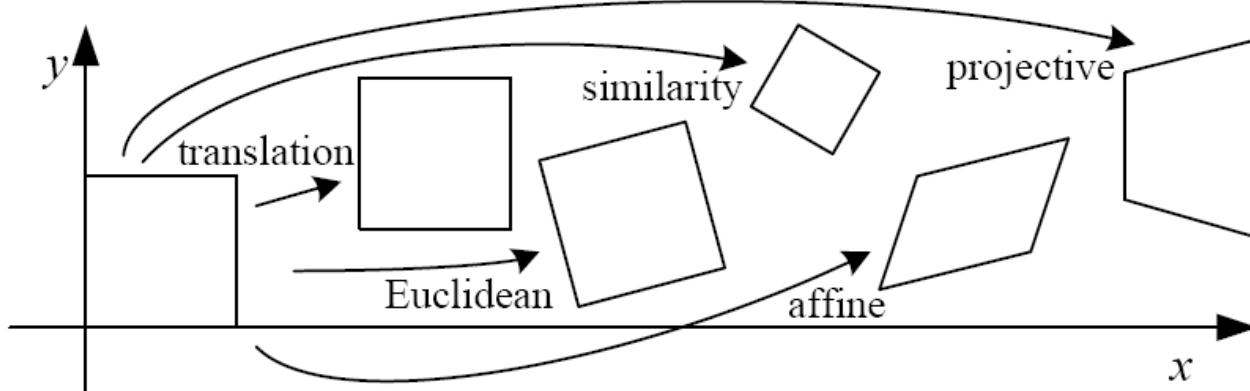
Linear Transformations as Change of Basis



$$\mathbf{p}' = \begin{bmatrix} u_x & v_x \\ u_y & v_y \end{bmatrix} \begin{bmatrix} 4 \\ 3 \end{bmatrix} = \begin{bmatrix} u_x & v_x \\ u_y & v_y \end{bmatrix} \mathbf{p}$$

- Any linear transformation is a basis!!!

2D Image Transformations

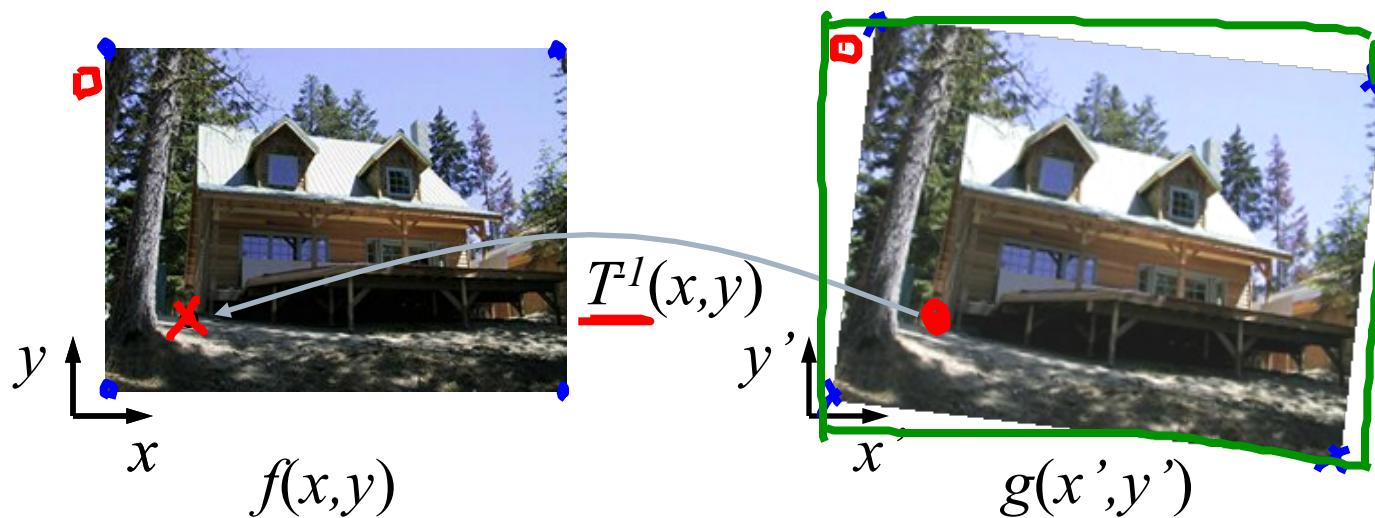


Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

These transformations are a nested set of groups

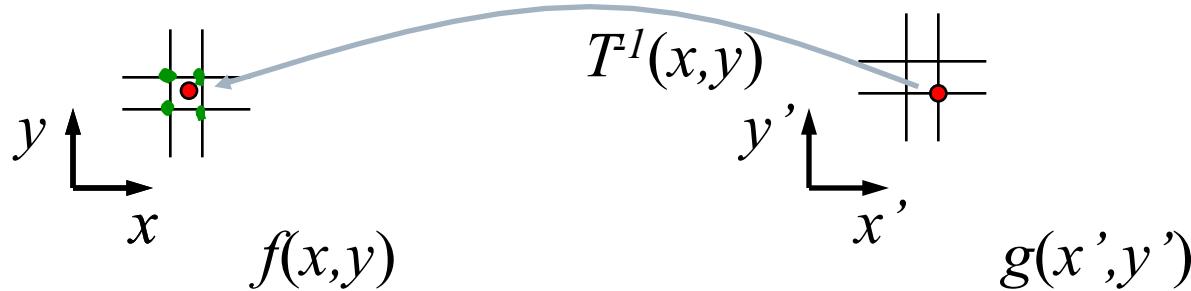
- Closed under composition and inverse is a member

Inverse Warping



- Get each pixel $g(x',y')$ from its corresponding location
 - $(x,y) = T^{-1}(x',y')$ in the first image
- Q: what if pixel comes from “between” two pixels?

Inverse Warping



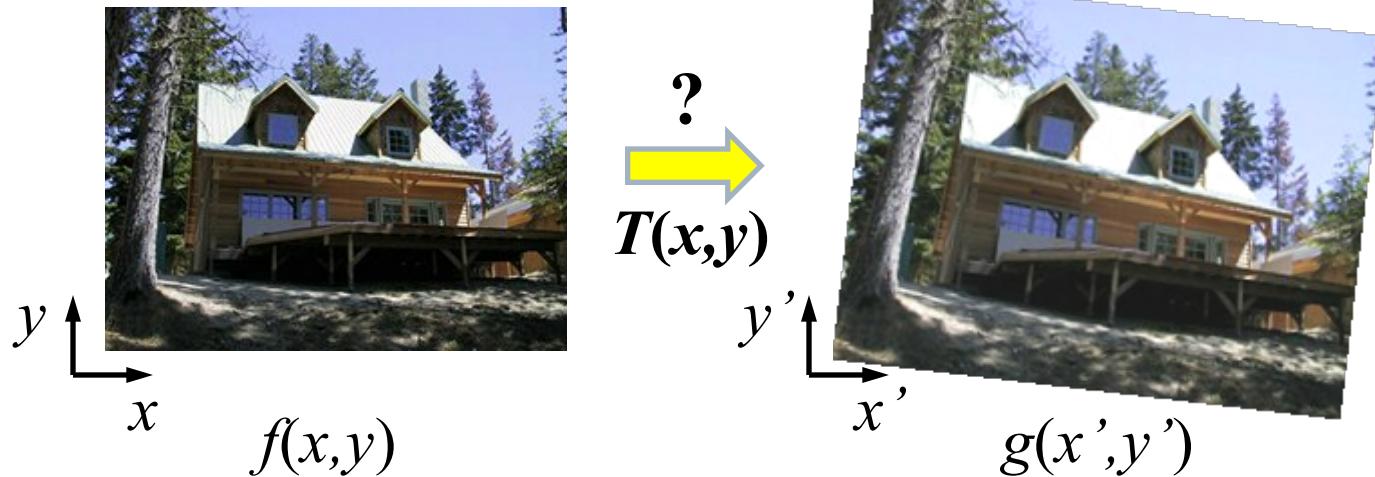
- Get each pixel $g(x',y')$ from its corresponding location
 - $(x,y) = T^{-1}(x',y')$ in the first image

Q: what if pixel comes from “between” two pixels?
A: *Interpolate* color value from neighbors

- nearest neighbor, bilinear, Gaussian, bicubic

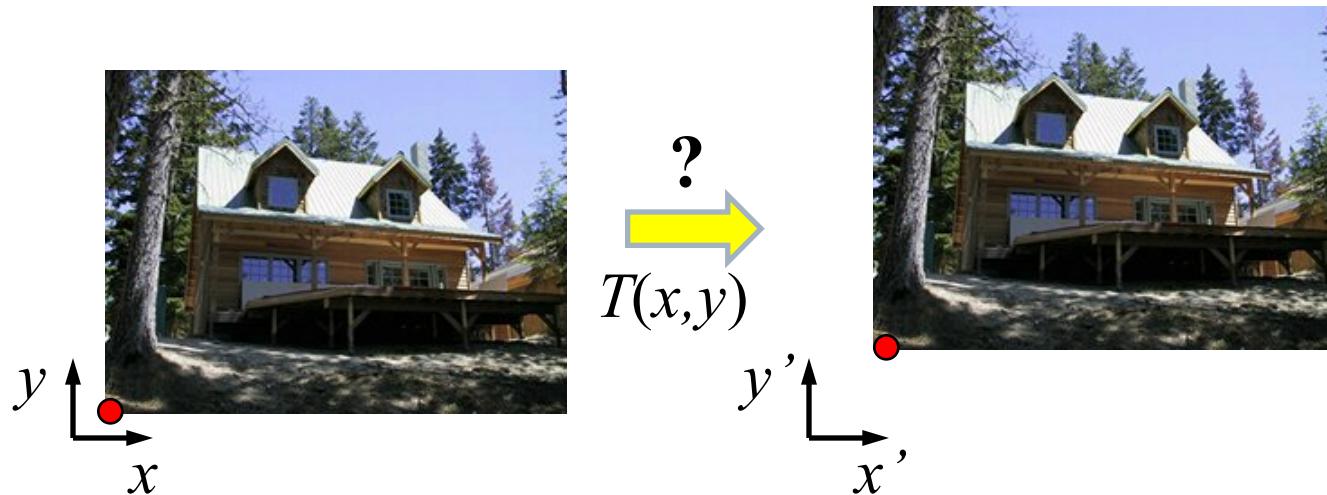
Image Morphing

Recovering Transformations



- What if we know f and g and want to recover the transform T ?
 - willing to let user provide correspondences
 - How many do we need?

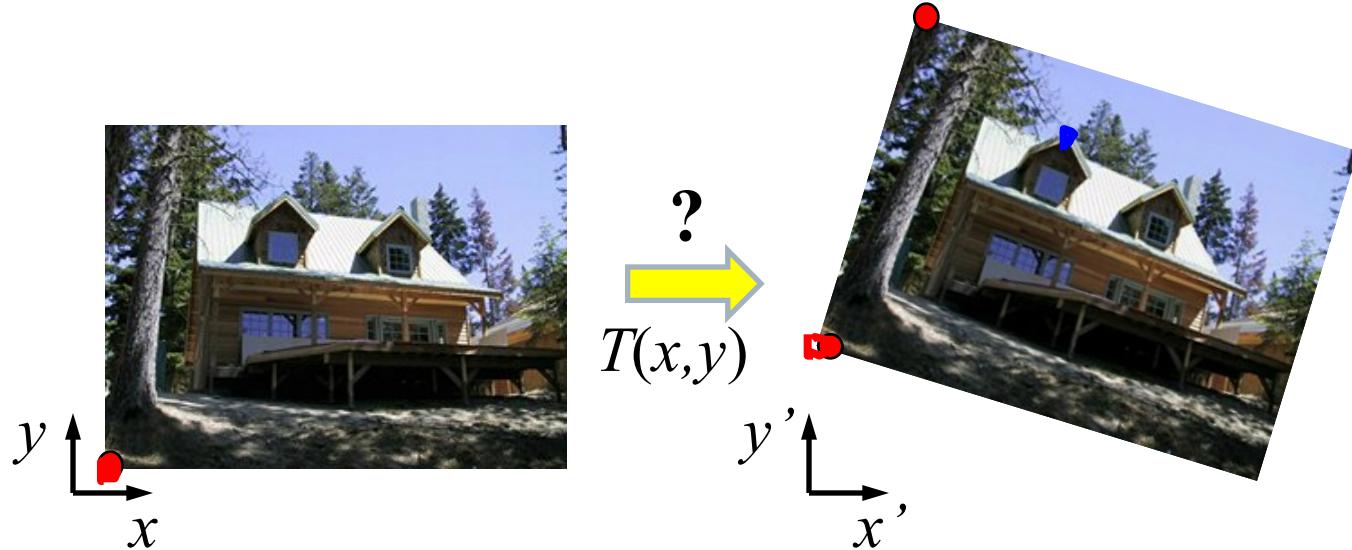
Translation: # correspondences?



- How many correspondences needed for translation?
- How many Degrees of Freedom?
- What is the transformation matrix?

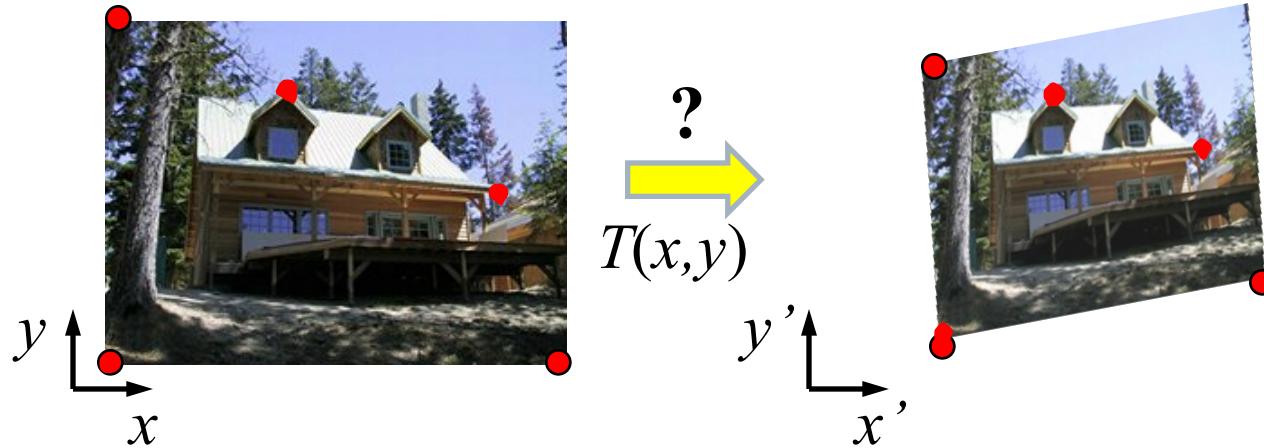
$$\mathbf{M} = \begin{bmatrix} 1 & 0 & p'_x - p_x \\ 0 & 1 & p'_y - p_y \\ 0 & 0 & 1 \end{bmatrix}$$

Euclidian: # correspondences?



- How many correspondences needed for translation+rotation?
- How many DOF?

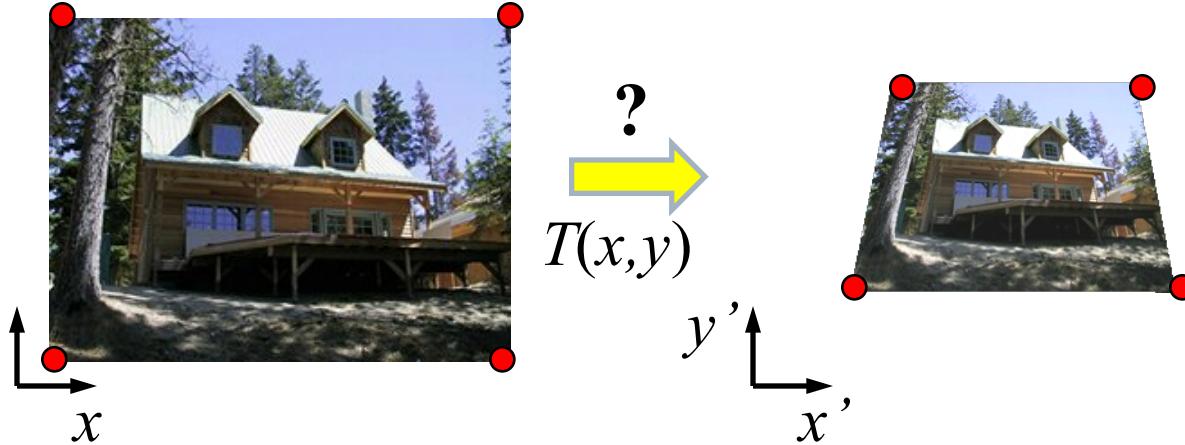
Affine: # correspondences?



- How many correspondences needed for affine?
- How many DOF?

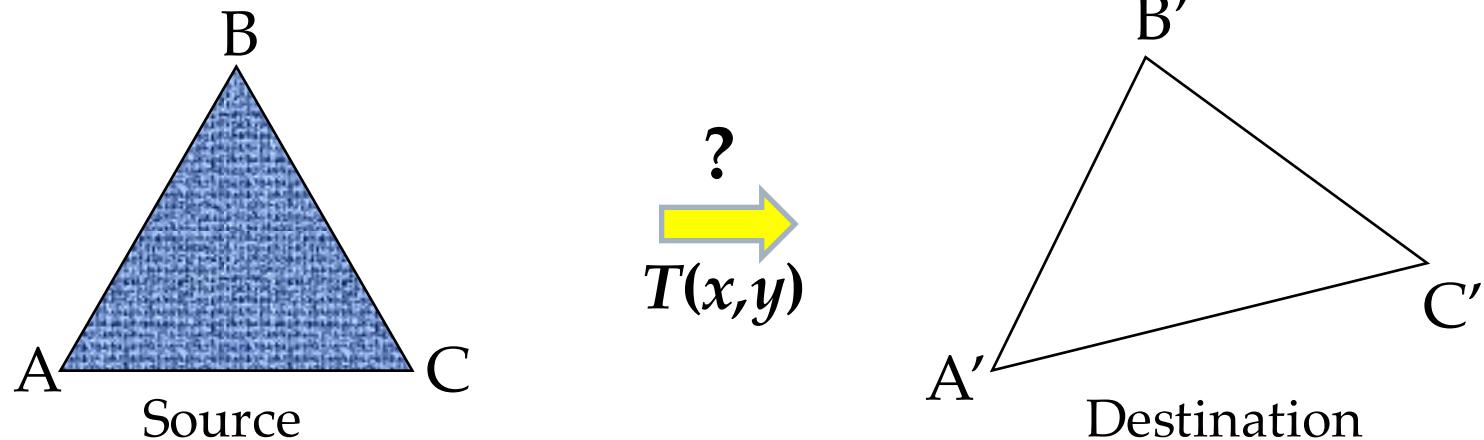


Projective: # correspondences?



- How many correspondences needed for projective?
- How many DOF?

Example: Warping Triangles

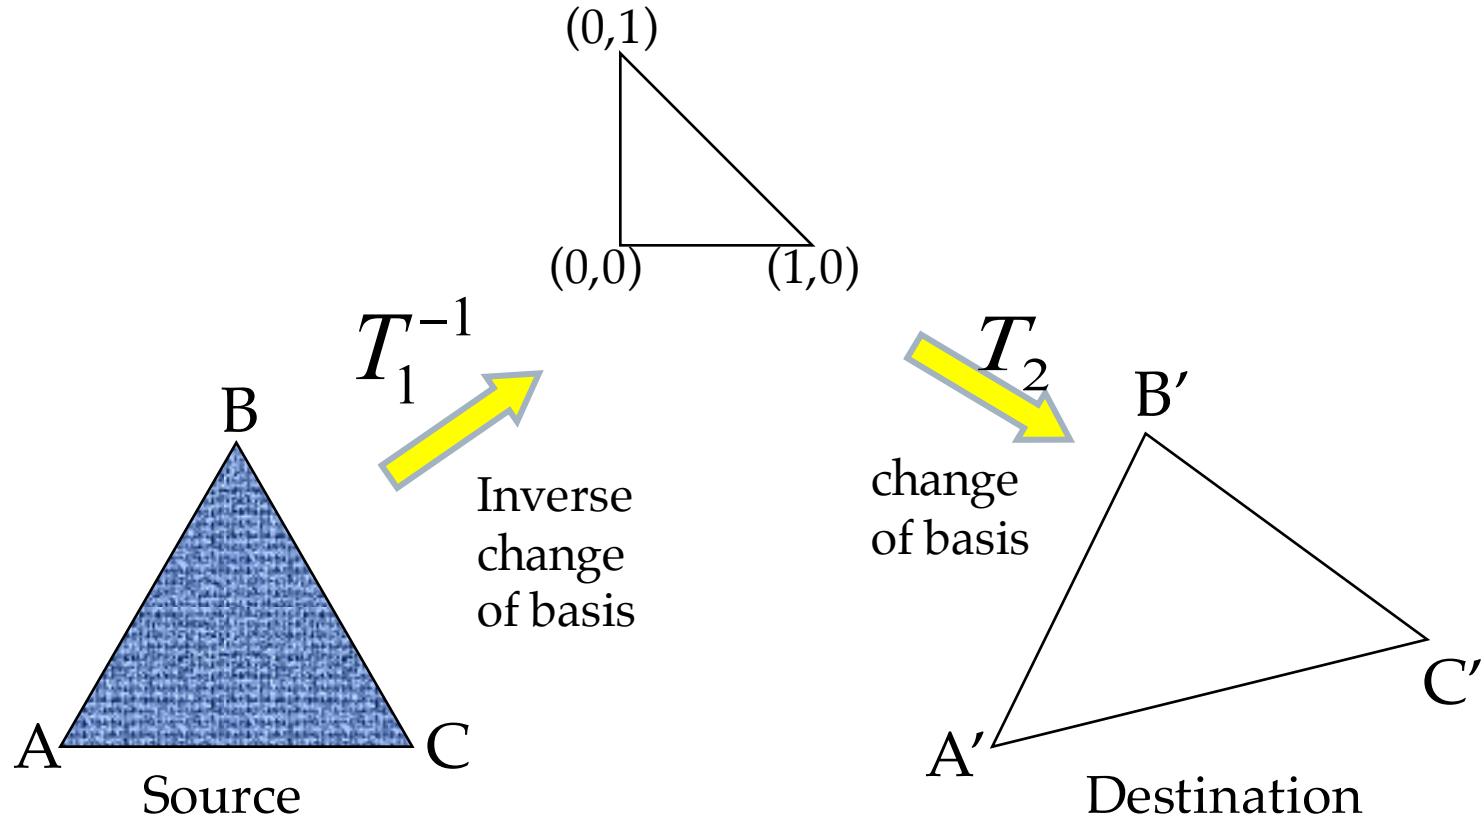


- Given two triangles: ABC and A'B'C' in 2D (12 numbers)
- Need to find transform T to transfer all pixels from one to the other.
- What kind of transformation is T?
- How can we compute the transformation matrix:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Least Squares!

HINT: Warping Triangles

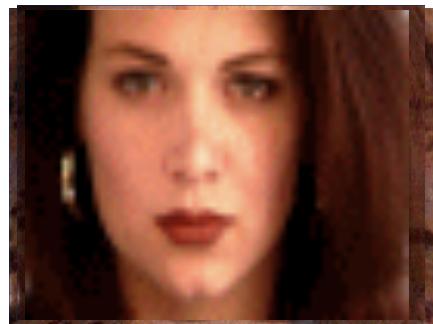


Don't forget to move the origin too!

- Very useful for Project 2...

Image Morphing

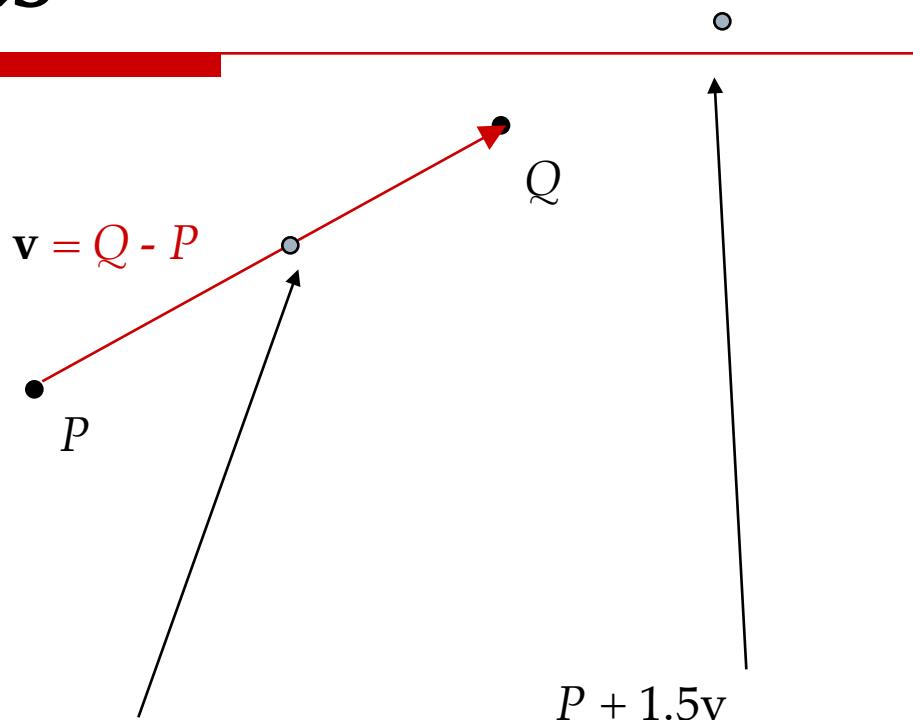
Morphing = Object Averaging



- The aim is to find “an average” between two objects
 - Not an average of two images of objects...
 - ...but an image of the average object!
 - How can we make a smooth transition in time?
 - Do a “weighted average” over time t
- How do we know what the average object looks like?
 - We don't have a clue!
 - But we can often fake something reasonable
 - Usually required user/artist input

Averaging Points

What's the average
of P and Q?



Linear Interpolation
(Affine Combination):
New point $aP + bQ$,
defined only when $a+b = 1$
So $aP+bQ = aP+(1-a)Q$

- P and Q can be anything:
 - points on a plane (2D) or in space (3D)
 - Colors in RGB or HSV (3D)
 - Whole images (m-by-n D)... etc.

Idea #1: Cross-Dissolve



- Interpolate whole images:
$$\text{Image}_{\text{halfway}} = (1-t) * \text{Image}_1 + t * \text{image}_2$$
- This is called **cross-dissolve** in film industry

- But what is the images are not aligned?

Idea #2: Align, then cross-dissolve



- Align first, then cross-dissolve
 - Alignment using global warp – picture still valid

Dog Averaging



- What to do?
 - Cross-dissolve doesn't work
 - Global alignment doesn't work
 - Cannot be done with a global transformation (e.g. affine)
 - Any ideas?
- Feature matching!
 - Nose to nose, tail to tail, etc.
 - This is a local (non-parametric) warp

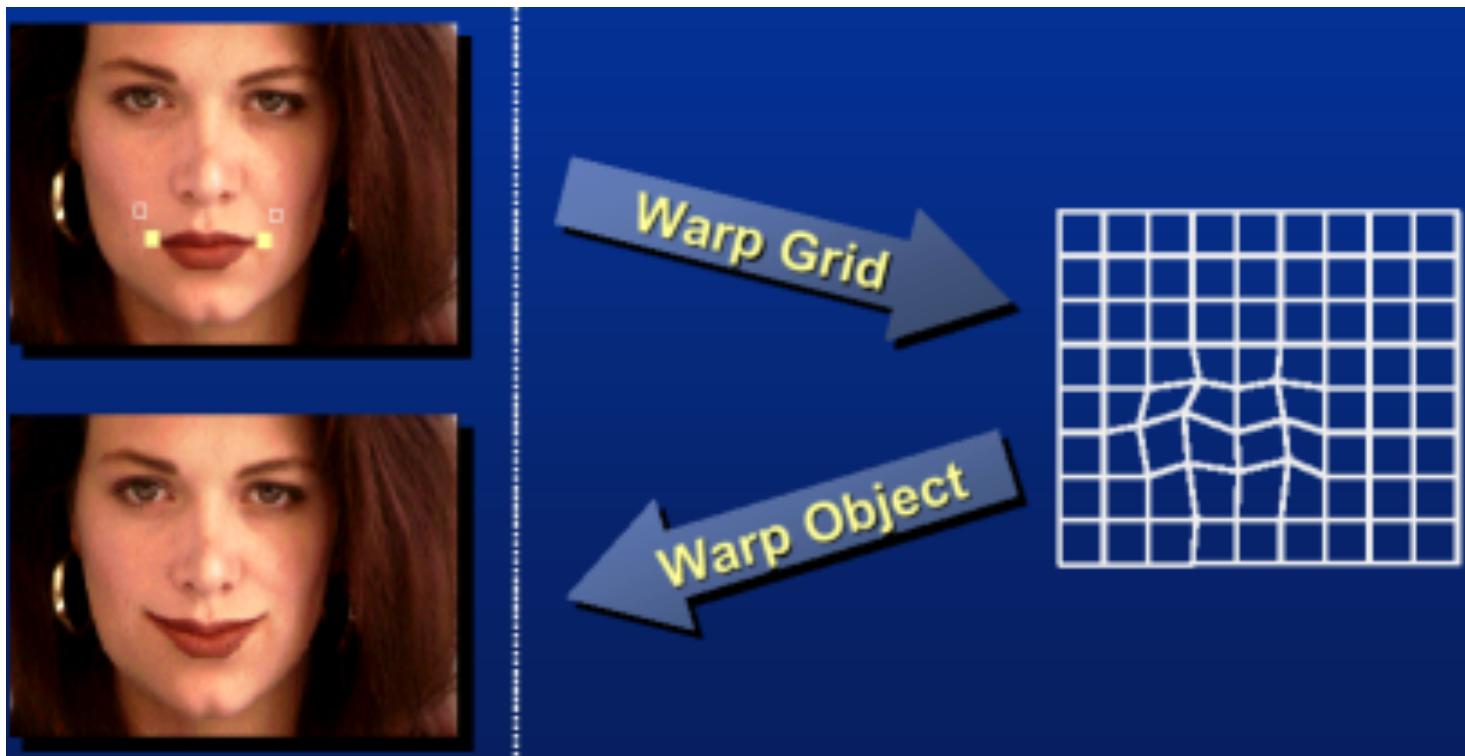
Local (non-parametric) Image Warping



- Need to specify a more detailed warp function
 - Global warps were functions of a few (2,4,8) parameters
 - Non-parametric warps $u(x,y)$ and $v(x,y)$ can be defined independently for every single location x,y !
 - Once we know vector field u,v we can easily warp each pixel (use backward warping with interpolation)

Image Warping – Non-parametric

- Move control points to specify a spline warp
- Spline produces a smooth vector field



Warp Specification - Dense

- How can we specify the warp?

Specify corresponding *spline control points*

- *interpolate* to a complete warping function



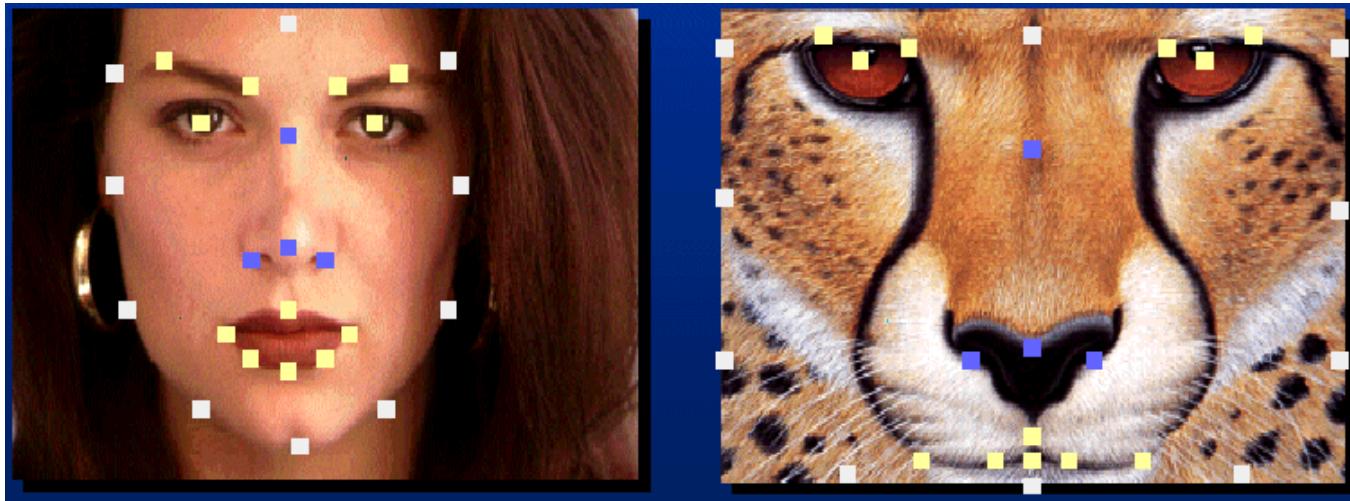
But we want to specify only a few points, not a grid

Warp Specification - Sparse

■ How can we specify the warp?

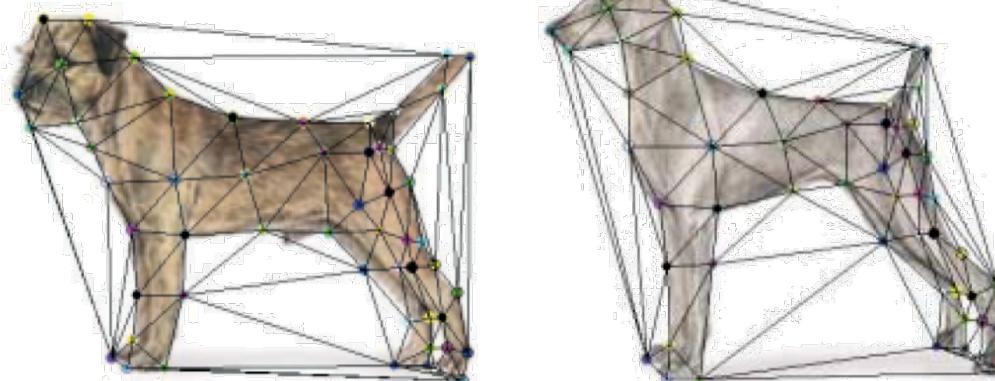
Specify corresponding *points*

- *interpolate* to a complete warping function
- How do we do it?



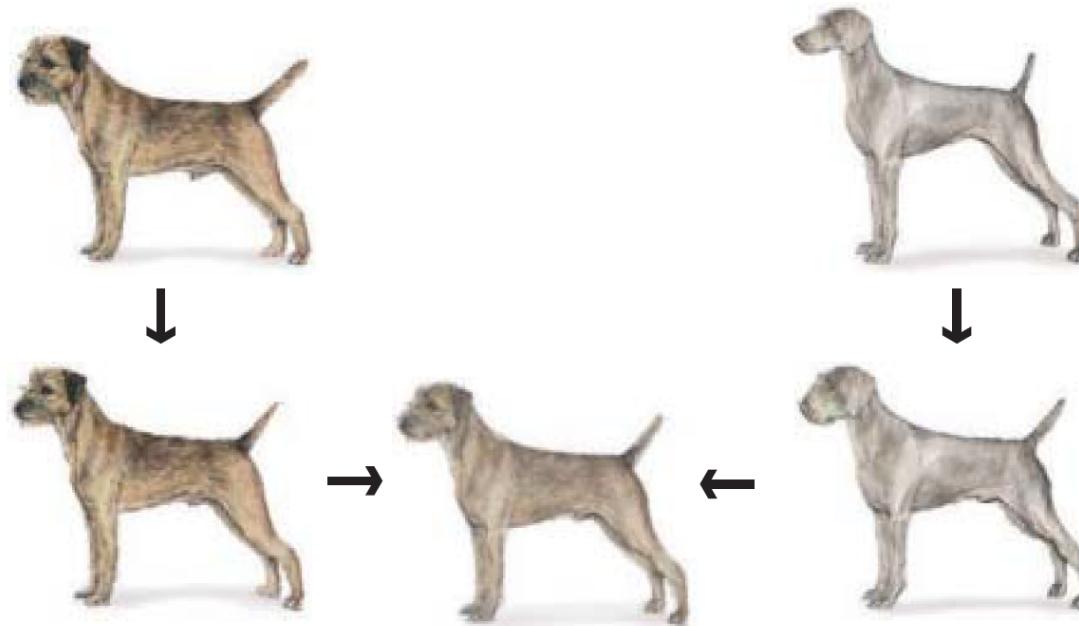
How do we go from feature points to pixels?

Triangular Mesh



1. Input correspondences at key feature points
2. Define a triangular mesh over the points
 - Same mesh in both images!
 - Now we have triangle-to-triangle correspondences
3. Warp each triangle separately from source to destination
 - How do we warp a triangle?
 - 3 points = affine warp!

Idea #3: Local warp, then cross-dissolve



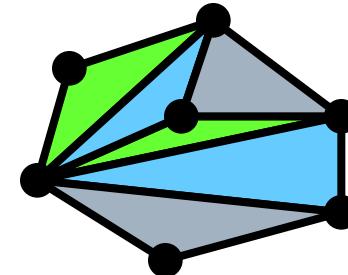
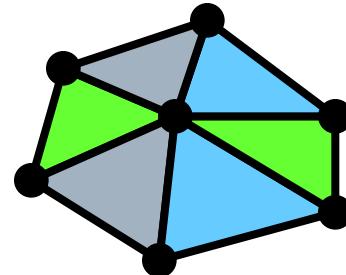
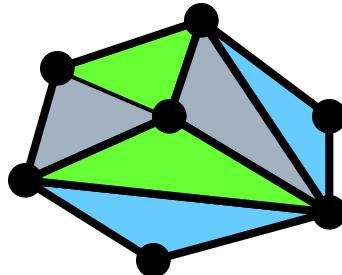
Morphing procedure:

for every triangle,

1. Find the average shape (the “mean dog” ☺)
 - local warping
2. Find the average color
 - Cross-dissolve the warped images

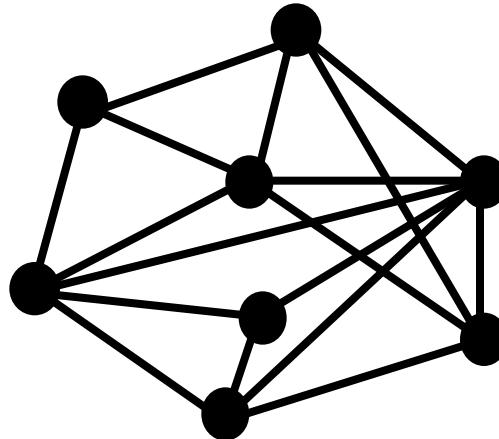
Triangulations

- A *triangulation* of set of points in the plane is a *partition* of the convex hull to triangles whose vertices are the points, and do not contain other points.
- There are an exponential number of triangulations of a point set.



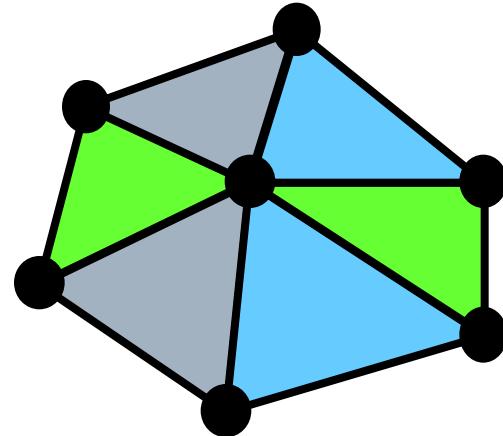
An $O(n^3)$ Triangulation Algorithm

- Repeat until impossible:
 - Select two sites.
 - If the edge connecting them does not intersect previous edges, keep it.

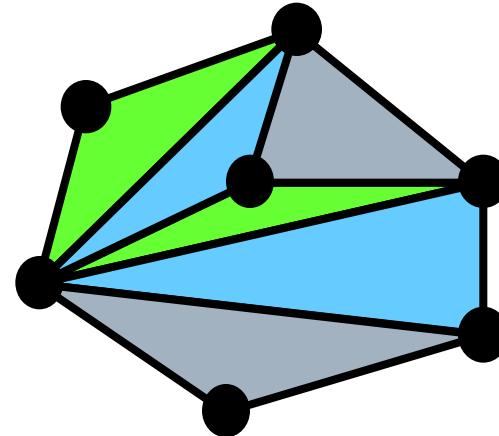


“Quality” Triangulations

- Let $a(T) = (a_1, a_2, \dots, a_{3t})$ be the vector of angles in the triangulation T in increasing order.
- A triangulation T_1 will be “better” than T_2 if $a(T_1) > a(T_2)$ lexicographically.
- The Delaunay triangulation is the “best”
 - Maximizes smallest angles



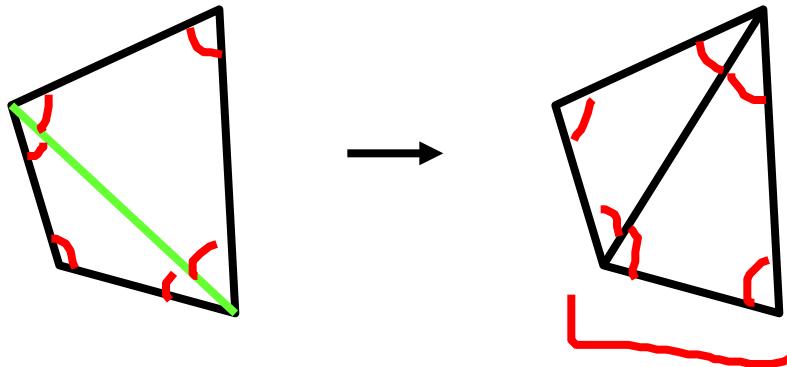
good



bad

Improving a Triangulation

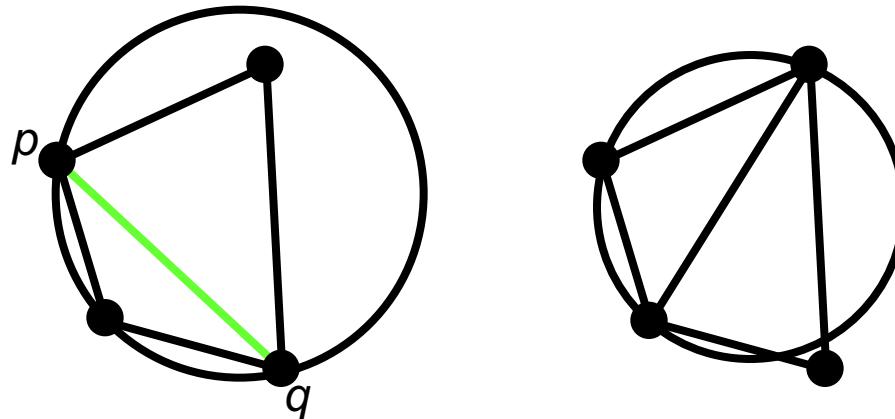
- In any convex quadrangle, an *edge flip* is possible. If this flip *improves* the triangulation locally, it also improves the global triangulation.



- If an edge flip improves the triangulation, the first edge is called *illegal*.

Illegal Edges

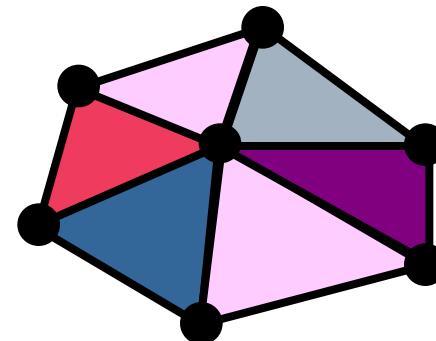
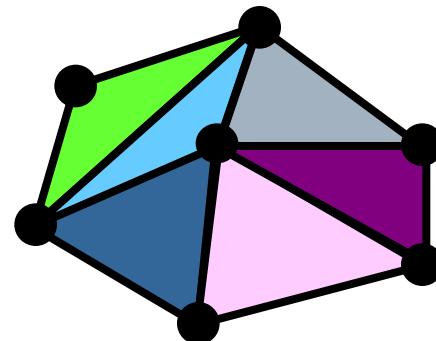
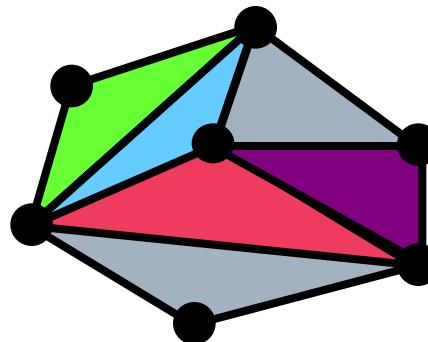
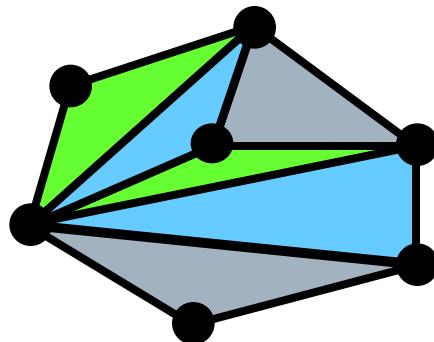
- **Lemma:** An edge pq is illegal iff one of its opposite vertices is inside the circle defined by the other three vertices.
- **Proof:** By Thales' theorem.



- **Theorem:** A Delaunay triangulation does not contain illegal edges.
- **Corollary:** A triangle is Delaunay iff the circle through its vertices is empty of other sites.
- **Corollary:** The Delaunay triangulation is not unique if more than three sites are co-circular.

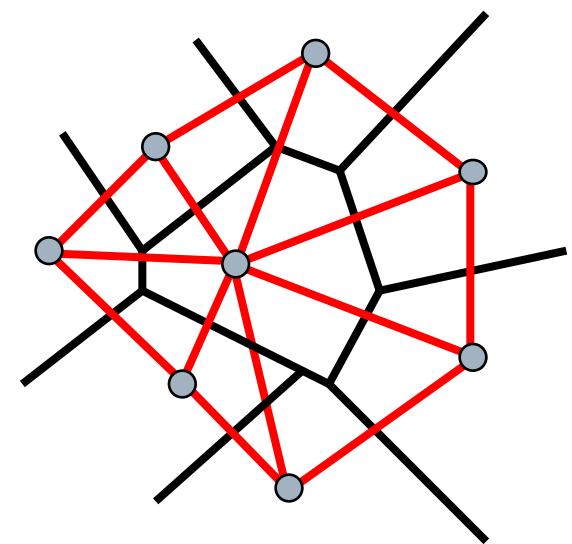
Naïve Delaunay Algorithm

- Start with an arbitrary triangulation. Flip any illegal edge until no more exist.
- Could take a long time to terminate.



Delaunay Triangulation by Duality

- General position assumption:
There are no four co-circular points.
- Draw the dual to the Voronoi diagram by connecting each two neighboring sites in the Voronoi diagram.
- **Corollary:** The DT may be constructed in $O(n \log n)$ time.
- This is what scipy's `delaunay` function uses.



Voronoi Planar Sweep

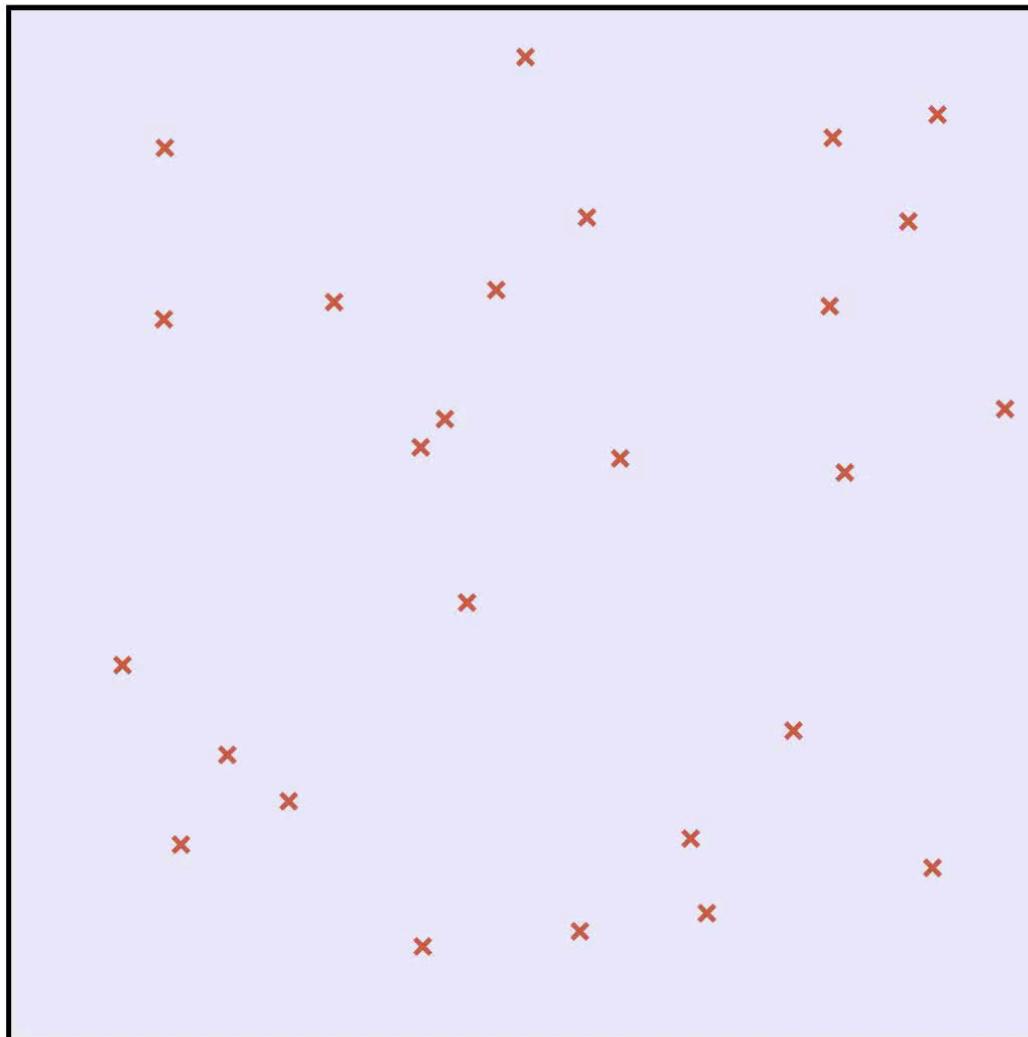


Image Morphing

- We know how to warp one image into the other, but how do we create a morphing sequence?
 1. Create an intermediate shape (by interpolation)
 2. Warp both images towards it
 3. Cross-dissolve the colors in the newly warped images

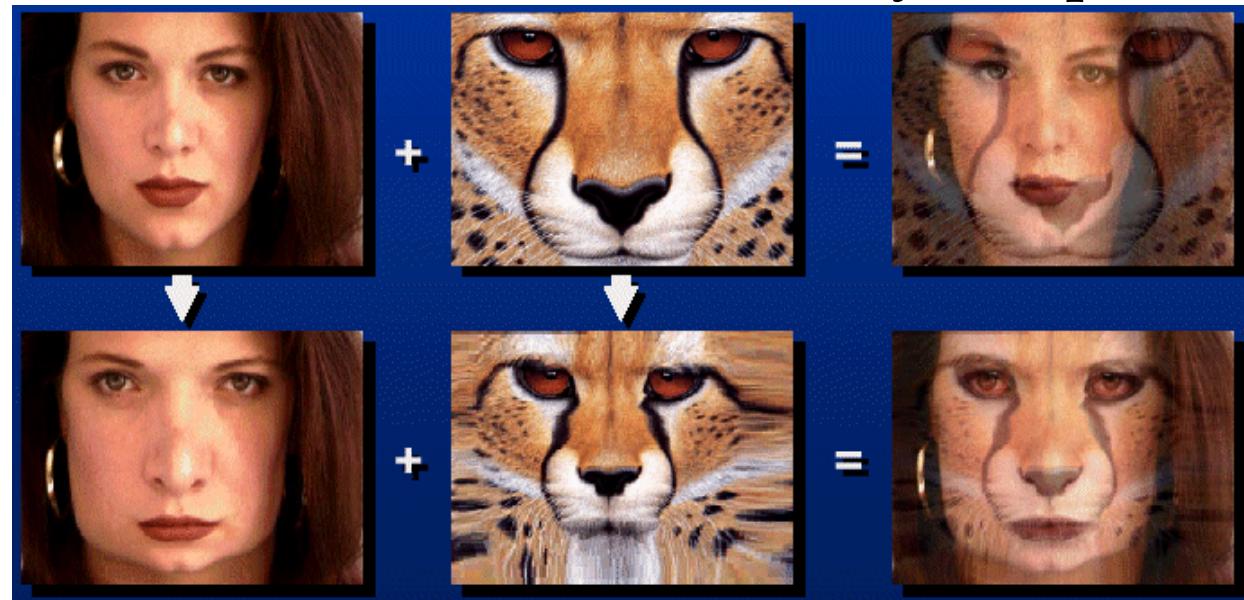


Image Morphing: Select Features

- 1) Generating correspondence points (by hand)

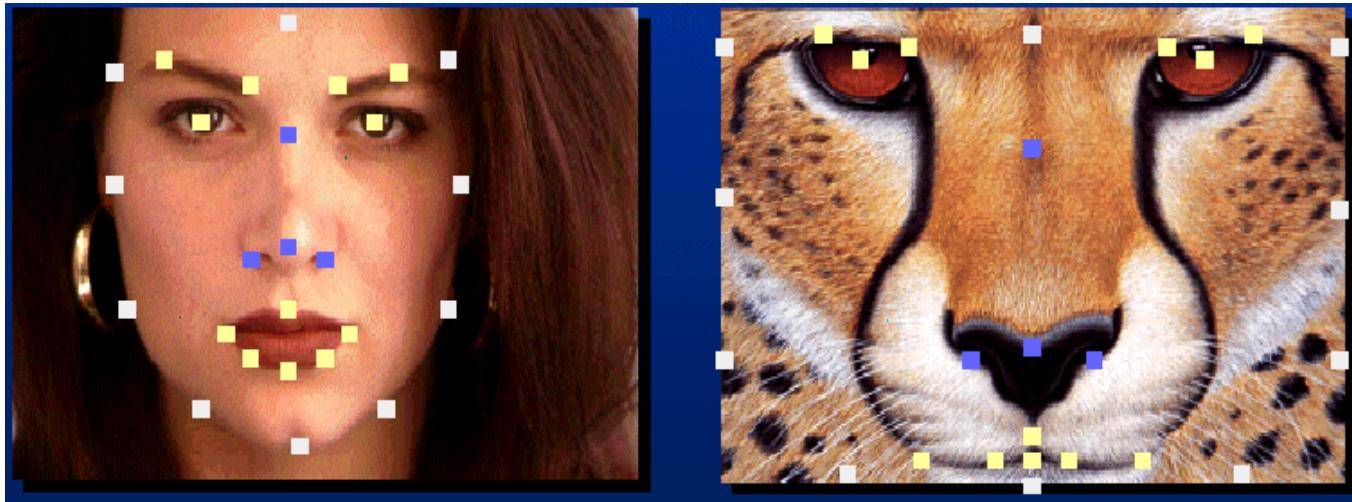


Image Morphing: Triangulation

- 2) Compute triangulation between the feature points
 - a) Use Delaunay triangulation (use library)
 - b) Make sure both images have the same triangulation!
 - a) **Use the intermediate shape to generate the triangulation (and use the fact you know correspondences of vertices).**

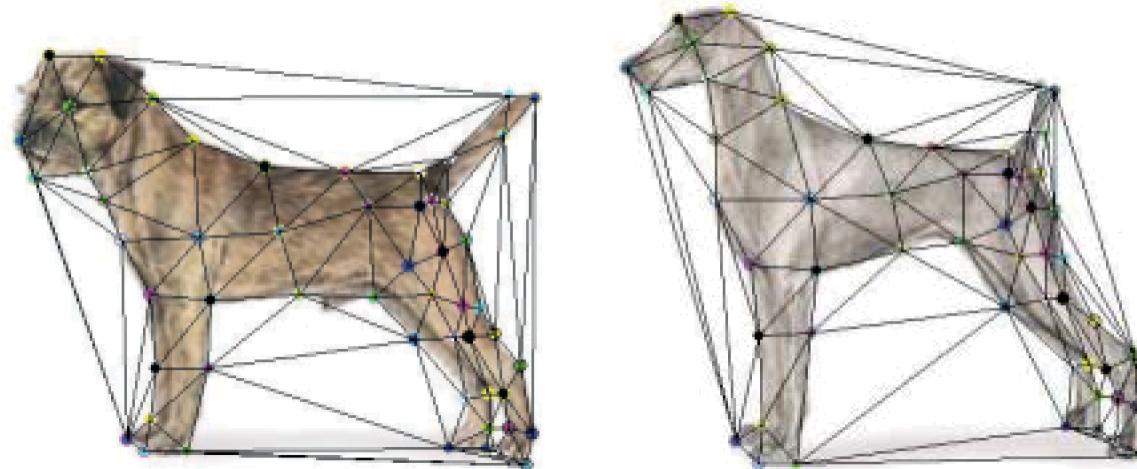
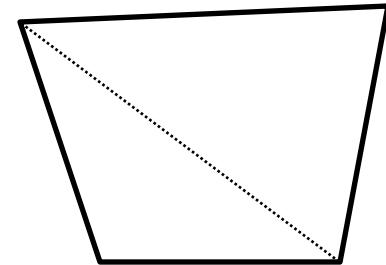
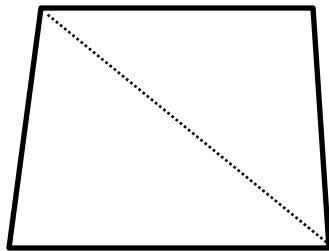
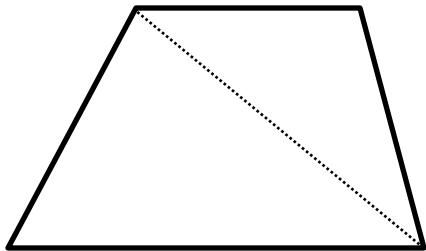


Image Morphing: Generate Intermediate Shape

2) Compute a weighted average shape

- Assume $t = [0,1]$
- Simple linear interpolation of each feature pair
- $(1-t)*p_1 + t*p_0$ for corresponding features p_0 and p_1



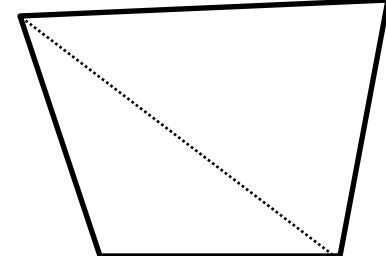
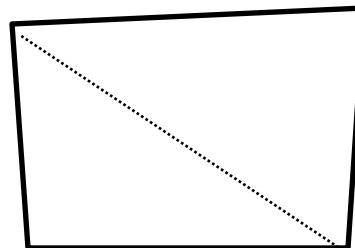
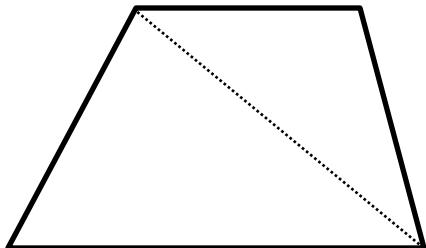
$t=0.4$



Image Morphing: Generate Intermediate Shape

2) Compute a weighted average shape

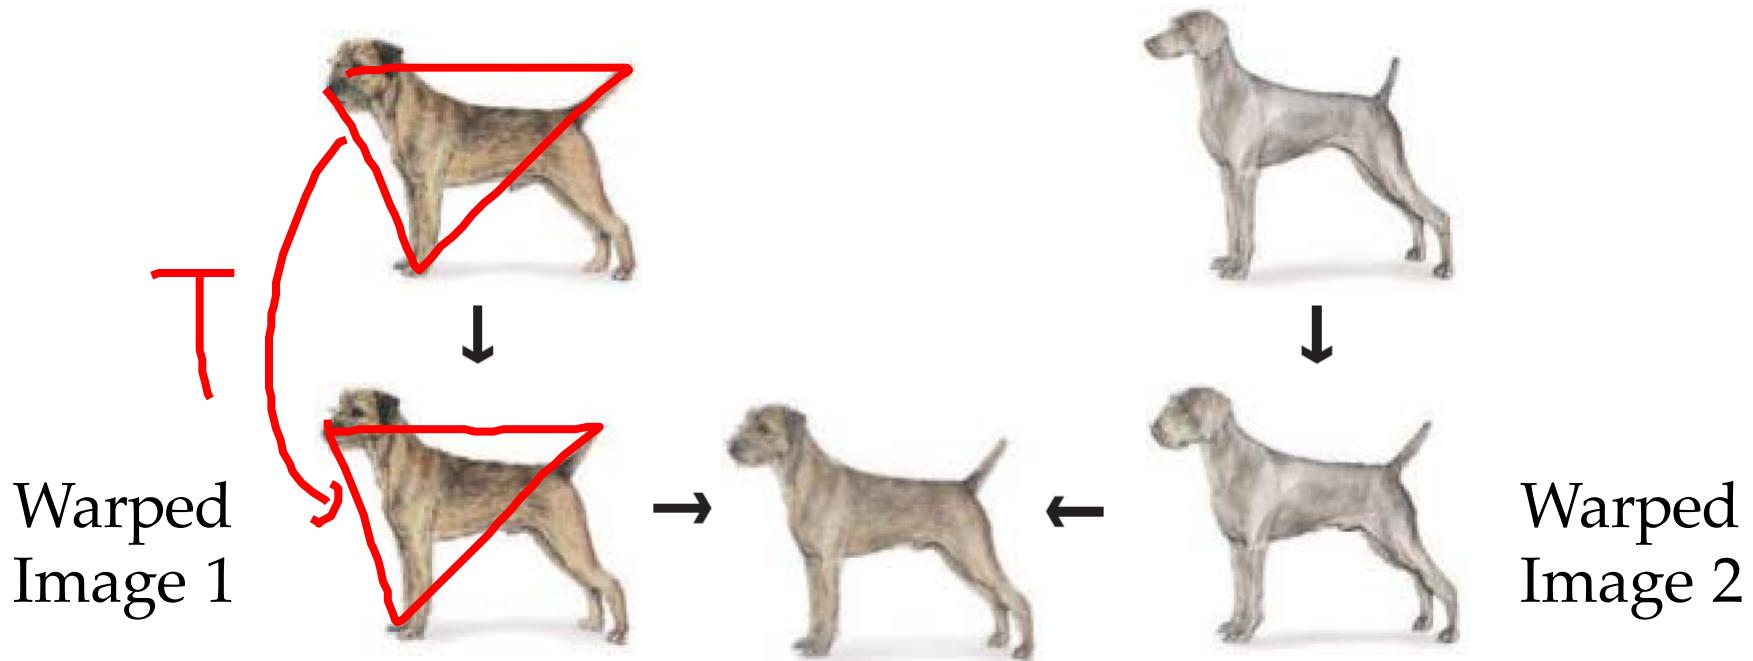
- Assume $t = [0,1]$
- Simple linear interpolation of each feature pair
- $(1-t)*p1+t*p0$ for corresponding features $p0$ and $p1$



$t=0.7$

In your implementation, compute the intermediate locations of the corresponding points (i.e., vertices) and then triangulate!!

Image Morphing: Generate Warped Image



- 3) Generate warped image 1 and 2:
source: original image 1, and 2
target: t -intermediate shape

Image Morphing: Generate Warped Image with Triangulation

Barycentric coordinates are invariant to affine warping!!

- Determine which triangle it is in
- Compute its barycentric coordinates
- Find equivalent point in equivalent triangle in original image, and copy over its intensity value (use those barycentric coordinates)

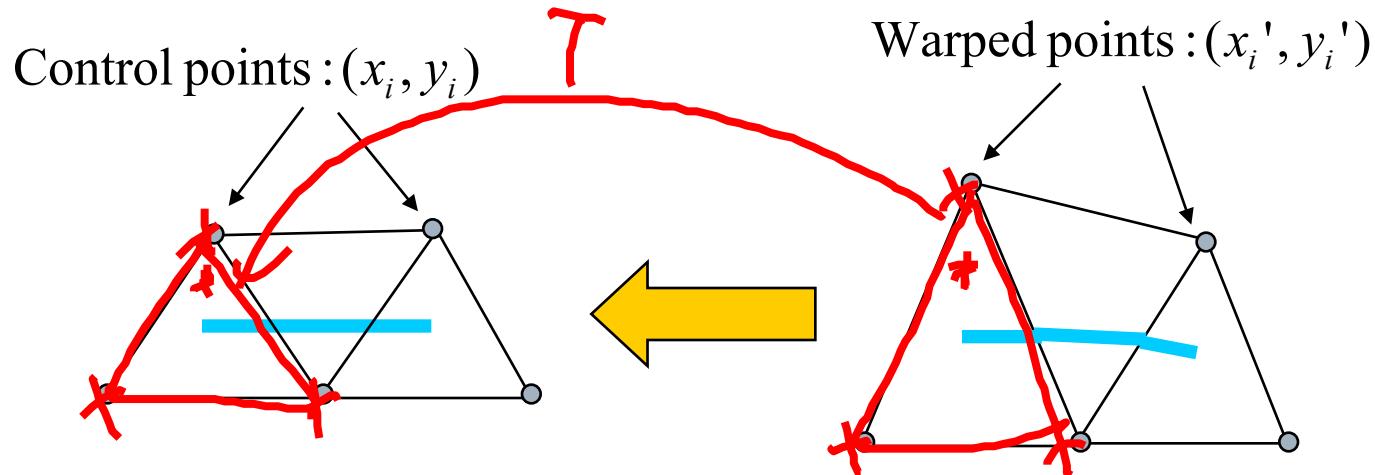
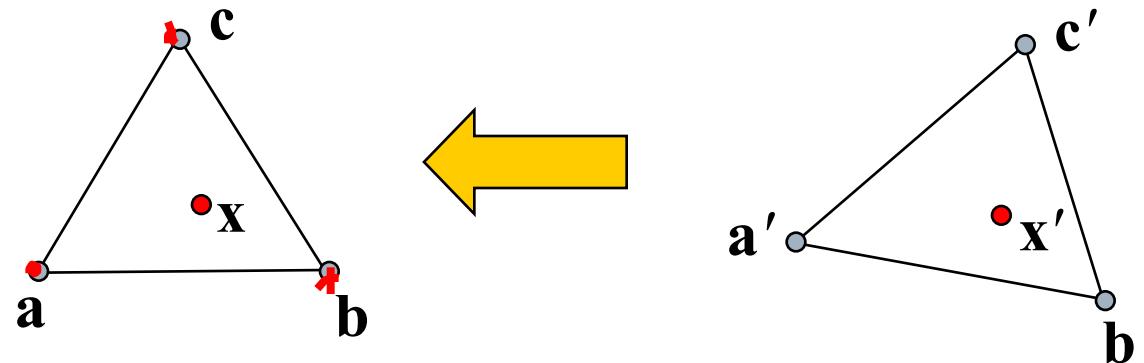


Image Morphing: Generate Warped Image with Triangulation

Barycentric coordinates are invariant to affine warping!!

- Determine which triangle it is in
- Compute its barycentric coordinates
- Find equivalent point in equivalent triangle in original image, and copy over its intensity value

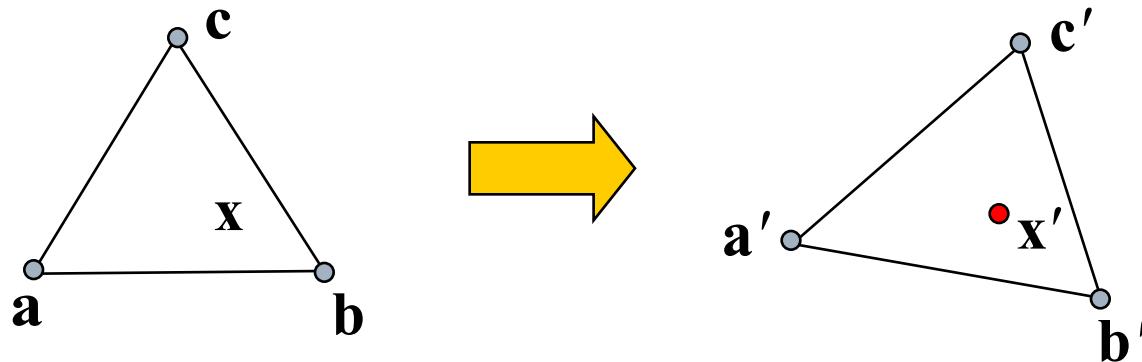


$$\mathbf{x} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$$

$$\alpha + \beta + \gamma = 1$$

$$\mathbf{x}' = \alpha\mathbf{a}' + \beta\mathbf{b}' + \gamma\mathbf{c}'$$

Barycentric Coordinates



$$\mathbf{x} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$$

$$\alpha + \beta + \gamma = 1$$

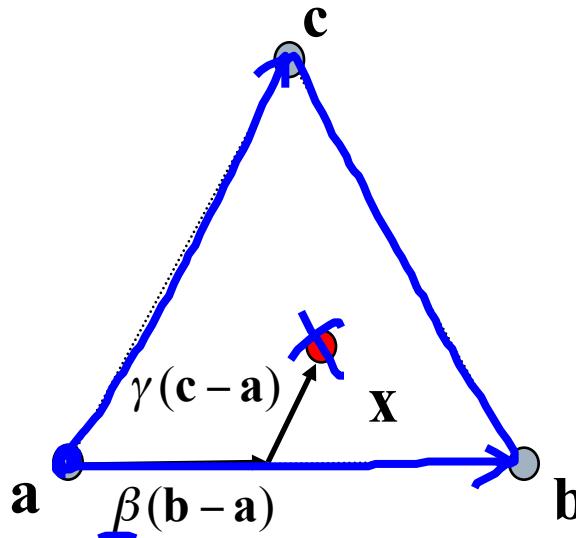
$$\mathbf{x}' = \alpha'\mathbf{a}' + \beta'\mathbf{b}' + \gamma'\mathbf{c}'$$

How do we know if a point
is inside of a triangle?

\mathbf{x} is inside the triangle if
 $0 \leq \alpha \leq 1$ and $0 \leq \beta \leq 1$

**In your implementation, use a scipy Delaunay method
to determine the triangle encasing that vertex!**

Barycentric Coordinates



$$\begin{aligned}\mathbf{x} &= \mathbf{a} + \beta(\mathbf{b} - \mathbf{a}) + \gamma(\mathbf{c} - \mathbf{a}) \\ &= (1 - \beta - \gamma)\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c} \\ &= \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}\end{aligned}$$

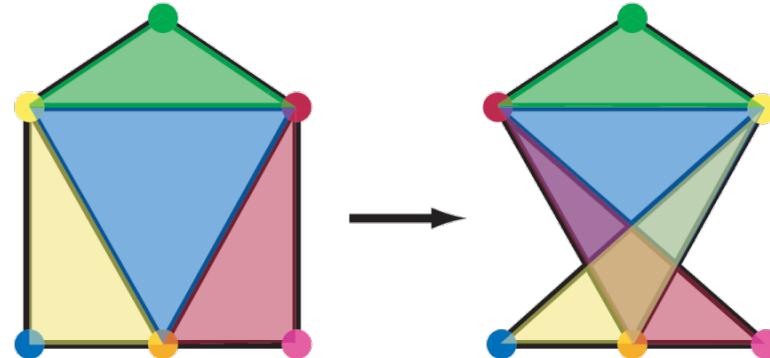
$$\mathbf{x} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$$

$$\alpha + \beta + \gamma = 1$$

$$\begin{pmatrix} x \\ y \\ 1 \end{pmatrix} = \begin{pmatrix} a_x & b_x & c_x \\ a_y & b_y & c_y \\ 1 & 1 & 1 \end{pmatrix} \begin{pmatrix} \alpha \\ \beta \\ \gamma \end{pmatrix}$$

Three linear equations in 3 unknowns

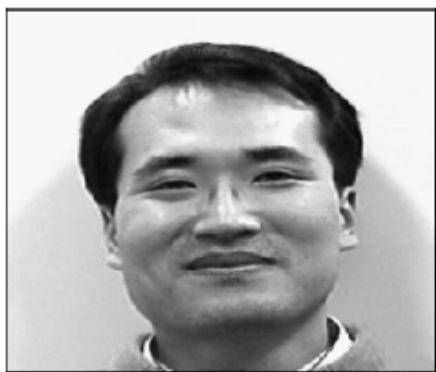
Other Issues



- Beware of folding
 - You are probably trying to do something 3D-ish
- Morphing can be generalized into 3D
 - If you have 3D data, that is!
- Extrapolation can sometimes produce interesting effects
 - Caricatures

Dynamic Scene

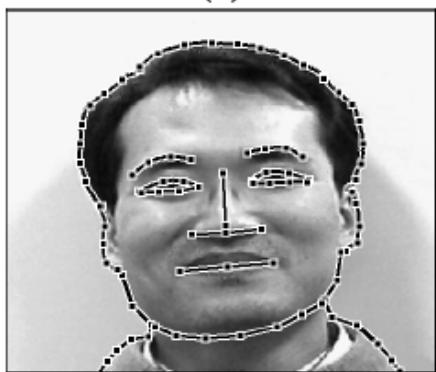




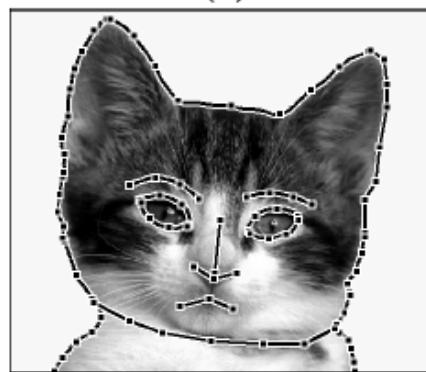
(a)



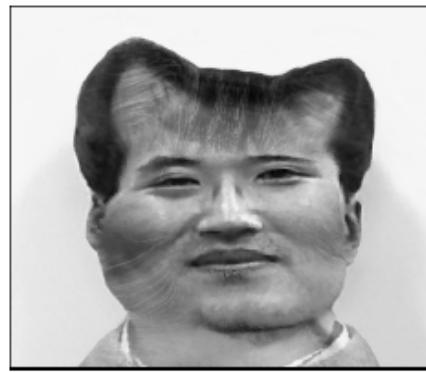
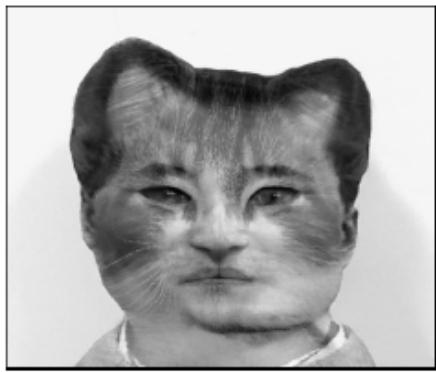
(b)

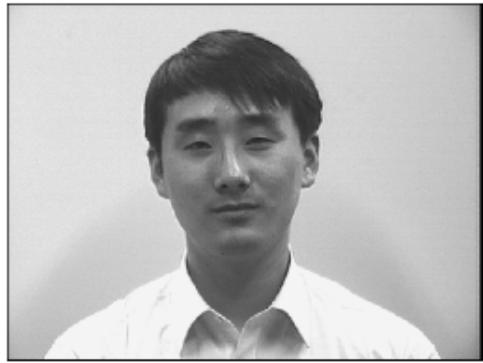


(c)

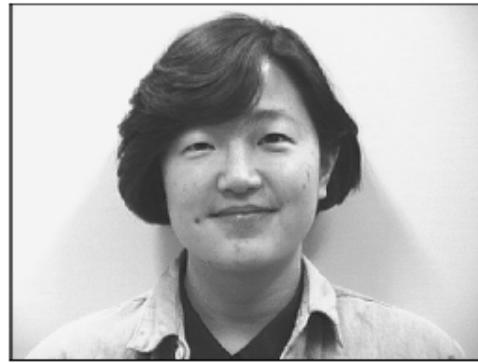


(d)

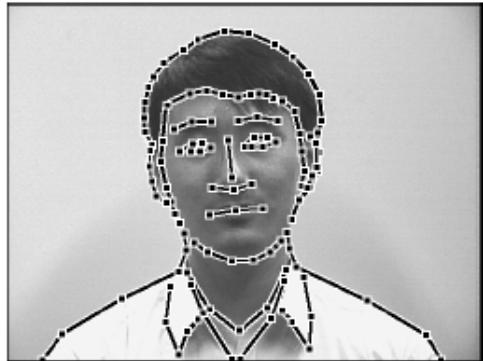




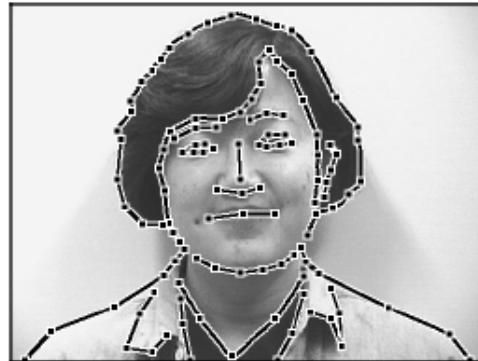
(a)



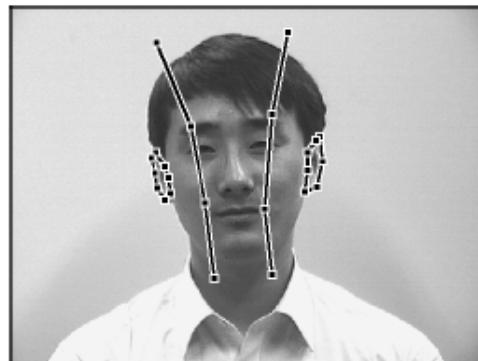
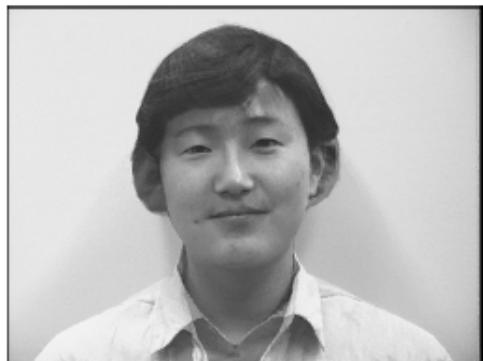
(b)



(c)



(d)



Project #3: Image Morphing

- Given two photos, produce a 60-frame morph animation



Recent Work in Face Morphing

- <https://www.youtube.com/watch?v=ohmajJTpNk>